

**ANTHROPIC**

# Research Engineer

San Francisco, CA / Remote-Friendly (Travel Required) / New York, NY / London, UK

Research / Full-Time / Hybrid

**Simplify****2 of 6 keywords**  
are present in your resumeBeta **33%** ×  
resume match

You want to build large scale ML systems from the ground up. You care about making safe, steerable, trustworthy systems. As a Research Engineer, you'll touch all parts of our code and infrastructure, whether that's making the cluster more reliable for our big jobs, improving throughput and efficiency, running and designing scientific experiments, or improving our dev tooling. You're excited to write code when you understand the research context and more broadly why it's important.

*Note: This is an "evergreen" role that we keep open on an ongoing basis. We receive many applications for this position, and you may not hear back from us directly if we do not currently have an open role on any of our teams that matches your skills and experience. We encourage you to apply despite this, as we are continually evaluating for top talent to join our team. You are also welcome to reapply as you gain more experience, but we suggest only reapplying once per year.*

*We may also put up separate, team-specific job postings. when In those cases, the teams will give preference to candidates who apply to the team-specific postings, so if you are interested in a specific team please make sure to check for team-specific job postings!*

## About Anthropic

Anthropic is an AI safety and research company that's working to build reliable, interpretable, and steerable AI systems. We want AI to be safe and beneficial for our customers and for society as a whole. Our interdisciplinary team has experience across ML, physics, policy, business and product.

### You may be a good fit if you:

- Have significant software engineering experience
- Are results-oriented, with a bias towards flexibility and impact
- Pick up slack, even if it goes outside your job description
- Enjoy pair programming (we love to pair!)
- Want to learn more about machine learning research
- Care about the societal impacts of your work

### Strong candidates may also have experience with:

- High performance, large-scale ML systems
- GPUs, Kubernetes, Pytorch, or OS internals
- Language modeling with transformers
- Reinforcement learning
- Large-scale ETL

### Representative projects:

- Optimizing the throughput of a new attention mechanism
- Comparing the compute efficiency of two Transformer variants
- Making a Wikipedia dataset in a format models can easily consume
- Scaling a distributed training job to thousands of GPUs
- Writing a design doc for fault tolerance strategies
- Creating an interactive visualization of attention between tokens in a language model

### Annual Salary:

- The expected salary range for this position is \$280k - \$600k USD.

### Logistics

**Location-based hybrid policy:** Currently, we expect all staff to be in our office at least 25% of the time.

**Deadline to apply:** None. Applications will be reviewed on a rolling basis.

**US visa sponsorship:** We do sponsor visas! However, we aren't able to successfully sponsor visas for every role and every candidate; operations roles are especially difficult to support. But if we make you an offer, we will make every effort to get you into the United States, and we retain an immigration lawyer to help with this.

LinkedIn tracking: #LI-DNI

**We encourage you to apply even if you do not believe you meet every single qualification.** Not all strong candidates will meet every single qualification as listed. Research shows that people who identify as being from underrepresented groups are more prone to experiencing imposter syndrome and doubting the strength of their candidacy, so we urge you not to exclude yourself prematurely and to submit an application if you're interested in this work. We think AI systems like the ones we're building have enormous social and ethical implications. We think this makes

representation even more important, and we strive to include a range of diverse perspectives on our team.

## **Compensation and Benefits\***

Anthropic's compensation package consists of three elements: salary, equity, and benefits. We are committed to pay fairness and aim for these three elements collectively to be highly competitive with market rates.

**Equity** - On top of this position's salary (listed above), equity will be a major component of the total compensation. We aim to offer higher-than-average equity compensation for a company of our size, and communicate equity amounts at the time of offer issuance.

**US Benefits** - The following benefits are for our US-based employees:

- Optional equity donation matching at a 3:1 ratio, up to 50% of your equity grant.
- Comprehensive health, dental, and vision insurance for you and all your dependents.
- 401(k) plan with 4% matching.
- 22 weeks of paid parental leave.
- Unlimited PTO – most staff take between 4-6 weeks each year, sometimes more!
- Stipends for education, home office improvements, commuting, and wellness.
- Fertility benefits via Carrot.
- Daily lunches and snacks in our office.
- Relocation support for those moving to the Bay Area.

**UK Benefits** - The following benefits are for our UK-based employees:

- Optional equity donation matching at a 3:1 ratio, up to 50% of your equity grant.
- Private health, dental, and vision insurance for you and your dependents.
- Pension contribution (matching 4% of your salary).
- 22 weeks of paid parental leave.
- Unlimited PTO – most staff take between 4-6 weeks each year, sometimes more!
- Health cash plan.
- Life insurance and income protection.
- Daily lunches and snacks in our office.

*\* This compensation and benefits information is based on Anthropic's good faith estimate for this position as of the date of publication and may be modified in the future. Employees based outside of the UK or US will receive a different benefits package. The level of pay within the range will depend on a variety of job-related factors, including where you place on our internal performance ladders, which is based on factors including past work experience, relevant education, and performance on our interviews or in a work trial.*


## How we're different

We believe that the highest-impact AI research will be big science. At Anthropic we work as a single cohesive team on just a few large-scale research efforts. And we value impact — advancing our long-term goals of steerable, trustworthy AI — rather than work on smaller and more specific puzzles. We view AI research as an empirical science, which has as much in common with physics and biology as with traditional efforts in computer science. We're an extremely collaborative group, and we host frequent research discussions to ensure that we are pursuing the highest-impact work at any given time. As such, we greatly value communication skills. We do not have boundaries between engineering and research, and we expect all of our technical staff to contribute to both as needed.

The easiest way to understand our research directions is to read our [recent research](#). This research continues many of the directions our team worked on prior to Anthropic, including: [GPT-3](#), [Circuit-Based Interpretability](#), [Multimodal Neurons](#), [Scaling Laws](#), [AI & Compute](#), [Concrete Problems in AI Safety](#), and [Learning from Human Preferences](#).

## Come work with us!

Anthropic is a public benefit corporation based in San Francisco. We offer competitive compensation and benefits, optional equity donation matching, generous vacation and parental leave, flexible working hours, and a lovely office space in which to collaborate with colleagues.

Jobs powered by  **LEVER**